Research Documentation: MCP Discovery & Kagent Integration

Research Documentation: MCP Discovery & Kagent Integration @

Executive Summary \mathscr{O}

This research document provides comprehensive analysis of building a Google Python ADK agent capable of autodiscovering MCP servers and converting configurations to Kagent format. The research covers protocol specifications, authentication frameworks, network discovery patterns, and production deployment architectures.

🔬 Research Methodology 🖉

Scope of Investigation @

- Google Python Agent Development Kit (ADK) capabilities and architecture patterns
- MCP (Model Context Protocol) v1.0+ specification and discovery mechanisms
- Network service discovery protocols (mDNS, SSDP, ARP scanning)
- OAuth 2.1 + PKCE authentication requirements for MCP
- Kagent schema requirements and CRD structures
- Python 3.12+ implementation patterns for production deployment

Key Research Questions @

- 1. How can Google ADK v1.0+ be leveraged for hierarchical agent architectures?
- 2. What are the technical requirements for MCP server autodiscovery?
- 3. How should OAuth 2.1 + PKCE authentication be implemented for automated systems?
- 4. What schema mapping strategies preserve semantic integrity between MCP and Kagent?
- 5. Which Python patterns optimize network discovery performance and reliability?

🟗 Google Python ADK Architecture Analysis 🛭

Framework Capabilities @

Core Agent Types

- LlmAgent: LLM-powered reasoning with model integration
- BaseAgent: Custom implementations with full control
- Workflow Agents: Sequential, Parallel, and Loop orchestration patterns

The ADK follows an event-driven, modular architecture supporting hierarchical multi-agent systems where parent agents coordinate sub-agents through structured event communication.

Production Integration Patterns

- Vertex AI Agent Engine: Fully managed runtime with auto-scaling capabilities
- Vertex AI Model Garden: Access to 200+ models including Gemini 2.0 Flash
- Cloud Service Integration: Direct BigQuery, AlloyDB, and Cloud Storage connectors
- Authentication: Application Default Credentials for development, Service Account integration for production

Security Framework

- OAuth 2.0, HTTPBearer, APIKey, and OpenID Connect schemes
- OpenAPI 3.0 compatibility with comprehensive validation

- Vertex Gemini Enterprise API for sandboxed code execution
- Callback systems (before_model_callback, before_tool_callback) for security controls

Implementation Recommendations @

Agent Hierarchy Design

Event Communication Pattern

- Type-safe inter-agent messaging using Pydantic models
- Circuit breaker patterns for external service calls
- · Structured logging with correlation IDs for distributed tracing

⊕ MCP Protocol Specification Analysis ℯ

Protocol Foundation \mathscr{O}

Transport Layer

- JSON-RPC 2.0: Core messaging protocol for all MCP communication
- STDIO: Local process execution for development and testing
- Server-Sent Events (SSE): Remote connections with real-time updates
- HTTP/WebSocket: API integration for production deployments

Capability Types

- Tools: Executable functions with JSON Schema parameter definitions
- Resources: Data sources with URI-based access patterns
- Prompts: Template systems for dynamic content generation
- Sampling: Advanced model interaction capabilities

Discovery Architecture @

Current Limitations

MCP v1.0+ relies primarily on manual configuration through JSON files. Network-based discovery requires custom implementation using standard protocols.

Proposed Discovery Methods

- mDNS/DNS-SD: Service announcement with _mcp._tcp.local service type
- SSDP/UPnP: Windows-compatible discovery for enterprise environments
- Custom Registry: Centralized service registration with health checking

Discovery Metadata Structure

```
MCPServiceRecord:
service_name: str
transport_type: Literal["stdio", "sse", "http", "websocket"]
endpoint_url: Optional[HttpUrl]
capabilities: List[MCPCapability]
auth_requirements: AuthenticationMetadata
version: str
```

🔐 OAuth 2.1 + PKCE Authentication Analysis 🛭

Protocol Requirements @

OAuth 2.1 Compliance

- PKCE (Proof Key for Code Exchange): Mandatory for all client types per RFC 7636
- Dynamic Client Registration: RFC 7591 for automated client setup
- Resource Indicators: RFC 8707 for proper token scoping
- Short-lived Tokens: 15-60 minute access token lifetime

Security Controls

- Exact redirect URI validation (no pattern matching)
- Audience validation with no token passthrough
- State parameter for CSRF protection
- Secure random generation for PKCE challenge

Implementation Architecture @

Authentication Flow

- 1. Discovery Phase: /.well-known/oauth-protected-resource endpoint discovery
- 2. Metadata Retrieval: Authorization server metadata via RFC 8414
- 3. Client Registration: Dynamic registration with capability declaration
- 4. Token Exchange: PKCE-enabled authorization code flow
- 5. Token Management: Automatic refresh with rotation

Credential Storage Strategy

- Encrypted at-rest storage for refresh tokens
- Memory-only storage for access tokens
- Automatic credential rotation on expiration
- Fallback authentication strategies for offline scenarios

Primary Technology: python-zeroconf *∂*

Technical Specifications

- Pure Python implementation with optional Cython extensions
- IPv4 and IPv6 support with no external dependencies
- · Async variants available through aiozeroconf wrapper
- Performance: 100+ service discovery within 30 seconds on typical networks

Implementation Pattern

```
class MCPDiscoveryService:
def __init__(self):
    self.zeroconf = AsyncZeroconf()
self.browser = AsyncServiceBrowser(
    zeroconf=self.zeroconf,
    type_="_mcp._tcp.local.",
```

```
handlers=[self._on_service_state_change]
8
           )
9
10
     async def _on_service_state_change(
11
         self,
12
         zeroconf: AsyncZeroconf,
13
         service_type: str,
14
        name: str,
        state_change: ServiceStateChange
16
       ) -> None:
17
          # Service discovery event handling
18
```

Complementary Discovery Protocols @

SSDP (Simple Service Discovery Protocol)

- UPnP compatibility for Windows environments
- UDP multicast discovery with XML device descriptions
- Essential for enterprise networks with mixed operating systems

ARP-based Network Scanning

- Direct Layer 2 enumeration for completeness
- Identifies potential MCP servers without service announcement
- Requires privilege escalation for raw socket access

Performance Optimization Strategies $\mathscr O$

Concurrent Discovery

- Asyncio semaphore-based concurrency control (default: 50 concurrent operations)
- Connection pooling with aiohttp.TCPConnector (100 total, 20 per host)
- Result caching with 30-60 second TTL for discovered services

Resource Management

- Memory usage under 50MB for typical home networks (100+ devices)
- CPU usage under 20% during active discovery phases
- Proper cleanup with asyncio context managers

Schema Mapping Architecture 🛭

Kagent CRD Structure Analysis @

Kubernetes Integration

- Custom Resource Definitions with apiVersion: kagent.dev/v1alpha1
- Agent and Tool resources with structured metadata
- Built-in support for tool categorization and risk assessment

Required Fields

```
apiVersion: kagent.dev/v1alpha1
kind: Tool
metadata:
name: mcp-tool-identifier
labels:
kagent.dev/category: "api|file_system|database|external_service|compute|utility"
```

```
kagent.dev/risk-level: "low|medium|high"

spec:

type: mcp # or builtin, http

description: "Tool description"

schema:

# JSON Schema for parameters

mcpConfig:

serverEndpoint: "https://mcp-server.example.com"

toolName: "original-mcp-tool-name"
```

MCP Tool Schema Structure ∅

JSON Schema Format

```
1 MCPTool:
2 name: str # Required
3
      description: str # Required
     inputSchema: Dict[str, Any] # JSON Schema Draft 7
     outputSchema: Optional[Dict[str, Any]]
       annotations: Optional[MCPAnnotations]
 6
7
8 MCPAnnotations:
9
      readOnlyHint: Optional[bool]
10
       destructiveHint: Optional[bool]
11
       idempotentHint: Optional[bool]
```

Bidirectional Conversion Strategy @

MCP → Kagent Conversion

- 1. Name Sanitization: Convert MCP names to Kubernetes-compliant identifiers
- 2. Schema Transformation: Wrap single-property schemas in object containers
- 3. Metadata Mapping: Transform MCP annotations to Kagent labels and categories
- 4. CRD Generation: Create properly formatted Kubernetes resources

Kagent → MCP Conversion

- 1. Schema Extraction: Extract JSON Schema from Kagent CRD spec
- 2. **Metadata Preservation**: Map Kagent categories to MCP annotations
- 3. Validation: Ensure MCP compliance with protocol specifications
- 4. Tool Registration: Format for MCP server tool registration

Quality Assurance

- 100% semantic preservation for core capabilities
- Zero data loss during bidirectional conversion
- JSON Schema Draft 7 validation for all transformations
- Property-based testing with Hypothesis for edge case discovery

🚀 Python 3.12+ Implementation Patterns 🛭

Performance Optimizations @

AsyncIO Enhancements

- 75% performance improvement in asyncio socket operations
- Eager task execution reducing scheduling overhead

· Optimized connection pooling with persistent connections

Modern Patterns

```
1 # Structured concurrency with TaskGroup
2 async def discover_services(self) -> List[MCPService]:
3
     async with asyncio.TaskGroup() as tg:
4
           mdns_task = tg.create_task(self._discover_mdns())
5
          ssdp_task = tg.create_task(self._discover_ssdp())
 6
           arp_task = tg.create_task(self._discover_arp())
7
8
       return self._merge_results(mdns_task.result(), ssdp_task.result(), arp_task.result())
9
10 # Robust timeout handling
11 async def authenticate_with_server(self, server: MCPServer) -> AccessToken:
       async with asyncio.timeout(30.0):
12
13
           return await self._oauth_flow(server)
```

Type Safety with Pydantic V2 ∅

Configuration Management

```
1 class MCPDiscoveryConfig(BaseModel):
2
     model_config = ConfigDict(
          str_strip_whitespace=True,
3
           validate_assignment=True,
 4
5
           use_enum_values=True
6
       )
7
8
       discovery_timeout: int = Field(default=30, ge=5, le=300)
9
       max_concurrent_discoveries: int = Field(default=50, ge=1, le=200)
10
       cache_ttl_seconds: int = Field(default=60, ge=10, le=3600)
11
12
       @field_validator('discovery_timeout')
13
       @classmethod
14
       def validate_timeout(cls, v: int) -> int:
15
          if v < 5:
16
               raise ValueError('Discovery timeout must be at least 5 seconds')
17
           return v
```

Package Management with UV @

Performance Benefits

- 10-100x faster installation compared to pip
- Rust-powered dependency resolution with conflict detection
- Unified replacement for pip, pip-tools, and virtualenv

Project Structure

```
1 [project]
2 name = "mcp-vacuum"
3 version = "0.1.0"
4 requires-python = ">=3.12"
5 dependencies = [
6     "google-adk>=1.0.0",
7     "pydantic>=2.0.0",
8     "asyncio-zeroconf>=0.3.0",
9     "httpx>=0.25.0",
```

```
10
       "typer>=0.9.0"
11 ]
12
13 [project.optional-dependencies]
14 \text{ dev} = [
15
     "pytest>=7.4.0",
    "pytest-asyncio>=0.21.0",
16
    "pytest-cov>=4.1.0",
17
     "ruff>=0.1.0",
19
       "mypy>=1.5.0"
20 ]
```

III Production Deployment Architecture @

Google Cloud Integration $\mathscr O$

Vertex AI Agent Engine

- Fully managed runtime with automatic scaling
- Built-in monitoring and observability
- Integrated security with IAM and VPC controls

Service Integration Strategy

- Cloud Storage: Configuration caching and service registry
- BigQuery: Discovery analytics and performance metrics
- Cloud Monitoring: Real-time alerting and dashboard visualization
- Cloud Logging: Structured log aggregation with correlation

Security Implementation @

Network Security

- VPC-native networking with private service connect
- Firewall rules limiting MCP server access to specific ports
- Network policy enforcement for discovered services

Identity and Access Management

- Service Account with minimal required permissions
- Workload Identity for Kubernetes deployment
- Credential rotation automation with Cloud KMS integration

Audit and Compliance

- Comprehensive audit logging for all discovery and conversion operations
- Data residency controls for multi-region deployments
- SOC 2 Type II compliance for enterprise deployments

📊 Performance and Reliability Requirements 🛭

Service Level Objectives (SLOs) @

Discovery Performance

- 95% of discovery operations complete within 30 seconds
- Support for 500+ concurrent service evaluations
- Memory usage remains below 200MB during peak operations

Authentication Reliability

- 99.9% OAuth flow success rate for valid configurations
- Token refresh operations complete within 5 seconds
- · Credential rotation with zero downtime

Schema Conversion Accuracy

- 100% semantic preservation for MCP core capabilities
- Conversion operations complete within 100ms for typical configurations
- Zero data loss during bidirectional transformations

Error Handling and Recovery @

Circuit Breaker Patterns

- Fail-fast for unreachable MCP servers
- Exponential backoff with jitter for retry operations
- Graceful degradation when discovery services are unavailable

Monitoring and Alerting

- Real-time metrics for discovery success rates
- Performance dashboards with latency percentiles
- Automated alerting for authentication failures

⊚ Implementation Recommendations *⊘*

Development Priorities @

Phase 1: Foundation (Weeks 1-2)

- 1. Project setup with UV and modern Python tooling
- 2. Core Pydantic models for MCP and Kagent schemas
- 3. Basic Google ADK agent framework implementation

Phase 2: Core Functionality (Weeks 3-5)

- 1. Multi-protocol network discovery implementation
- 2. OAuth 2.1 + PKCE authentication system
- 3. MCP JSON-RPC 2.0 client with full transport support

Phase 3: Integration (Weeks 6-8)

- 1. Bidirectional schema conversion engine
- 2. ADK agent orchestration and lifecycle management
- 3. Comprehensive testing suite with real-world scenarios

Phase 4: Production (Weeks 9-12)

- 1. Google Cloud deployment automation
- 2. Monitoring, alerting, and observability
- 3. Performance optimization and security hardening

Technical Debt Mitigation @

Code Quality Standards

• 95%+ test coverage for core logic components

- Type safety with mypy strict mode
- Comprehensive documentation with architectural decision records

Operational Excellence

- Infrastructure as Code with Terraform
- CI/CD pipeline with automated security scanning
- Disaster recovery procedures with documented runbooks

📚 References and Further Reading 🛭

Primary Sources @

- Google ADK Documentation: https://google.github.io/adk-docs/
- MCP Specification: https://modelcontextprotocol.io/specification/
- OAuth 2.1 RFC: https://datatracker.ietf.org/doc/draft-ietf-oauth-v2-1/
- PKCE RFC 7636: https://tools.ietf.org/html/rfc7636
- Dynamic Client Registration RFC 7591: https://tools.ietf.org/html/rfc7591

Technical Standards @

- JSON-RPC 2.0: https://www.jsonrpc.org/specification
- mDNS RFC 6763: https://tools.ietf.org/html/rfc6763
- DNS-SD RFC 6763: https://tools.ietf.org/html/rfc6763
- SSDP UPnP Specification: http://upnp.org/specs/arch/UPnP-arch-DeviceArchitecture-v1.0.pdf

Implementation Guides @

- Python AsyncIO Best Practices: https://docs.python.org/3/library/asyncio.html
- Pydantic V2 Documentation: https://docs.pydantic.dev/latest/
- pytest-asyncio Patterns: https://pytest-asyncio.readthedocs.io/

Research conducted: June 24, 2025

Document version: 1.0.0 Next review: July 24, 2025